

Looking for Trends

300958 Social Web Analysis

Week 11 Lab Solutions

- Complete the following R-code to write a function to compute Twitter's χ^2 statistic for trends.

```
TwitterChi = function(tweetsBefore, tweetsNow, N) {

  return((tweetsNow - tweetsBefore)^2/tweetsBefore
    + ((N - tweetsNow) - (N - tweetsBefore))^2/(N - tweetsBefore))

}
```

Using the matrix:

	Before	Now	Sample Size
Number of Tweets	51	52	1000

Assignment Project Exam Help

<https://powcoder.com>

We can convert it into a matrix:

```
> Y = matrix(0, 2, 2)
> Y[, 1] = X[1:2]
> Y[, 2] = X[3] - X[1:2]
> colnames(Y) = c("Topic", "No Topic")
> rownames(Y) = c("Before", "After")
```

Add WeChat powcoder

We can stretch it into the long form:

```
> stretchTable = function(tab, variableNames) {
+   tabx = rep(rownames(tab), rowSums(tab))
+   l = ncol(tab)
+   m = nrow(tab)
+   cn = colnames(tab)
+   taby = c()
+   for (a in 1:m) {
+     for (b in 1:l) {
+       taby = c(taby, rep(cn[b], tab[a, b]))
+     }
+   }
+ }
```

```

+
+   d = data.frame(x = tabx, y = taby)
+   colnames(d) = variableNames
+   return(d)
+ }
>
> Z = stretchTable(Y, c("Time", "Topic"))

```

And examine the top few items in the stretched list.

```
> head(Z)
```

```

      Time Topic
1 Before Topic
2 Before Topic
3 Before Topic
4 Before Topic
5 Before Topic
6 Before Topic

```

Assignment Project Exam Help

Then tabulate it to confirm that we have not made any mistakes:

```
> table(Z)
```

	Topic	
Time	No Topic	Topic
After	948	52
Before	949	51

<https://powcoder.com>

Add WeChat powcoder

Compute the χ^2 randomisation distribution.

```

> ## define the functions to use
> expectedIndependent = function(X) {
+   n = sum(X)
+   p = rowSums(X)/sum(X)
+   q = colSums(X)/sum(X)
+   return(p %o% q * n) # outer product creates table
+ }
>
> chiSquaredStatistic = function(X, E) {
+   return(sum((X - E)^2/E))
+ }
>


```

```

> ## compute the expected table
> E = expectedIndependent(table(Z)) # compute expected counts if independent
>
> ## compute the randomisation distribution
> x2dist = replicate(1000, { # compute 1000 randomised chi-squared statistics
+   timeShuffle = sample(Z$Time)
+   topicShuffle = sample(Z$Topic)
+   Yindep = table(timeShuffle, topicShuffle)
+   chiSquaredStatistic(Yindep, E)
+ })
>
> hist(x2dist)

```

Histogram of x2dist



plot of chunk unnamed-chunk-7

Compute the χ^2 statistic for the original data.

```

> x2 = chiSquaredStatistic(table(Z), E)

```

Compute the p value of the test.

```
> ## pval is the proportion of x2dist that is greater than x2
> pval = mean(x2dist > x2)
> print(pval)
```

```
[1] 0.829
```

- Using R's χ^2 test:

```
> chisq.test(Y, simulate.p.value = TRUE)
```

Pearson's Chi-squared test with simulated p-value (based on 2000 replicates)

```
data: Y
X-squared = 0.010236, df = NA, p-value = 1
```

- Linear regression

```
> count = c(135, 145, 138, 142, 145, 103, 128, 144,
+ 149, 130, 107, 106, 83, 117, 123, 104, 116, 127, 75, 128, 136,
+ 145, 120, 127, 110, 109, 122, 136, 125, 143, 177, 142, 134, 153,
+ 173, 151, 174, 168, 158, 156, 128, 156, 110, 133, 132)
> day = c(210,
+ 211, 212, 213, 214, 215, 216, 217, 218, 219, 220, 221, 222, 223,
+ 224, 225, 226, 227, 228, 229, 230, 231, 232, 233, 234, 235, 236,
+ 237, 238, 239, 240, 241, 242, 243, 244, 245, 246, 247, 248, 249,
+ 250, 251, 252, 253, 254)
>
> m = lm(sqrt(count) ~ day)
> summary(m)
```

Call:

```
lm(formula = sqrt(count) ~ day)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.6112	-0.5232	0.0051	0.5408	1.5870

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.80411	2.36352	1.186	0.241973

```

day          0.03714    0.01017    3.651 0.000703 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.8862 on 43 degrees of freedom
Multiple R-squared:  0.2366,    Adjusted R-squared:  0.2189
F-statistic: 13.33 on 1 and 43 DF,  p-value: 0.0007034

```

The above summary shows the intercept 2.80411, and the gradient 0.03714. The p -value for the gradient (0.000703), shows that the population gradient is not likely to be zero.

- Moving averages using a loop

The completed function is:

```

> windowWeights = function(m) {
+   if( m%%2 ) {
+     ## m is odd
+     w = rep(1,m)/m
+   } else {
+     ## m is even
+     w = c(0.5, rep(1,m-1), 0.5)/m
+   }
+ }
>
> moving.average = function(x,m) {
+   ## compute window weights
+   w = windowWeights(m)
+   j = floor(m/2)
+   offsets = (-j):j
+   n = length(x)
+   res = rep(NA, n)
+   ## slide window and apply window weights to obtain averages
+   for(i in (j+1):(n-j)) {
+     res[i] = sum(w*x[i+offsets])
+   }
+   return(res)
+ }

```

- Make sure you understand this function and use it to compute the 7 point moving average of the tweet count data. Remember to square root first.

```

> trend = moving.average(sqrt(count), 7)
> print(trend)

```

```
[1]      NA      NA      NA 11.03508 11.08952 11.11308 11.09440
[8] 11.12934 11.13629 10.95317 10.88216 10.75223 10.46530 10.37510
[15] 10.50729 10.27367 10.58842 10.70917 10.84504 10.95310 11.02440
[22] 10.91278 11.16707 11.12873 11.12873 11.00570 11.14910 11.43977
[29] 11.64381 11.80603 11.99517 12.20817 12.36644 12.54253 12.49358
[36] 12.58693 12.71752 12.56672 12.47201 12.40686 12.16996 11.95962
[43]      NA      NA      NA
```

The following function computes a seasonal component by averaging over all observations that are m apart.

```
seasonal = function(x ,m) {
  ## extend the data so it is a multiple of m long
  tmp = c(x, rep(NA, m - length(x)%m))
  ## convert to a matrix
  mat = matrix(tmp, nrow=m)
  ## Calculate the row means to get the seasonal component (excluding missing entries)
  seas = rowMeans(mat, na.rm=TRUE)
  seas = seas - mean(seas)
  return(seas)
}
```


<https://powcoder.com>

- Make sure you understand this function and use it to compute the seasonal component for the tweet count data. Remember to subtract the trend first.

```
> centred = sqrt(count) - trend # remove trend
> zcentred = centred[!is.na(centred)] # remove NAs at boundaries
> s = seasonal(zcentred, 7) # compute seasonal component
> print(s)
```

```
[1] -0.2961236 -0.7975847 -0.4794883  0.4469257  0.3593245  0.2767446
[7]  0.4902018
```

```
> plot(day, sqrt(count))
> lines(day, trend, col=2)
> lines(day[!is.na(centred)], trend[!is.na(centred)] + rep_len(s, length.out =
sum(!is.na(centred))))
```



Assignment Project Exam Help

<https://powcoder.com>

Data split into trend and seasonal components.

Add WeChat powcoder